



# DATA MODELING *in the* 24 ½ CENTURY

Amanda Moran  
Developer Advocate DataStax

[velocityconf.com/ca](http://velocityconf.com/ca)  
#VelocityConf



# The Future Is Here!



Actually it's been here...

NoSQL has been around since the early 2000's



# What Will We Talk About?

- Why is everyone still talking about NoSql?
- Why is everyone talking about Apache Cassandra?
- What makes Data Modeling for Apache Cassandra different?
- How to fix common mistakes

# Who is Amanda?



Santa Clara  
University



Apache  
Trafodion



UDACITY

TERADATA ASTER

DATASTAX



O'REILLY  
Velocity

[velocityconf.com/ca](http://velocityconf.com/ca)  
#VelocityConf

# Why NoSQL?

- Built to solve issues with Relational Databases
- All NoSQL Databases have in common
  - Made for Big Data
  - Horizontally scalable
  - High Throughput (low latency)
  - High Availability
- Issues of yesterday, we have in spades today!



# Why Apache Cassandra?

- Many NoSQL Databases
- 10 years old -- it's survived!
- Open Source with a strong community
- Built for the cloud
  - Cloud agnostic
- Masterless Architecture -- True High Availability



# Why Apache Cassandra?

- Back-end of some of the most popular apps
  - Uber, Netflix, Hulu, Twitter, Apple
- Contributors to the code
  - DataStax, Apple, Twitter, Facebook



# The Basics


- First developed by Facebook
- Top-level Apache Foundation project in 2010
- ***Distributed***, decentralized database
- Elastic scalability
  - Add/remove nodes with no downtime






# The Basics

- High performance
  - Very fast -- low latency
- High availability / fault tolerant
  - No single point of failure



# Let's Talk about the Big Topics


- Distributed Systems
- Replication
- Elastic Scalability
- High Availability
- Performance/Latency
  - Read Path
  - Write Path



Note: Don't forget this is just a brief intro!


# Distributed System

- Masterless Architecture
- All nodes have the same job
- All Clients can Connect to All Nodes
- All nodes ready for Read and Write
- Not all data on all nodes




# Replication

- Data must be copied to other nodes
  - High Availability
- Replication Factor (RF) is set by the user
  - 1 - # of Nodes
- Asynchronously replicated
  - Automatic
  - Peer-to-peer communication



# Token Range and The Ring






# Elastic Scalability

- Nodes ++ = Performance ++
- Scale up or down with no downtime
  - Not even a restart!
- Reads and Writes both scale

## Scale-Up Linearity




# High Availability

- Master node allows for high availability
  - No single point of failure
- Replication allows nodes to fail and data to still be available

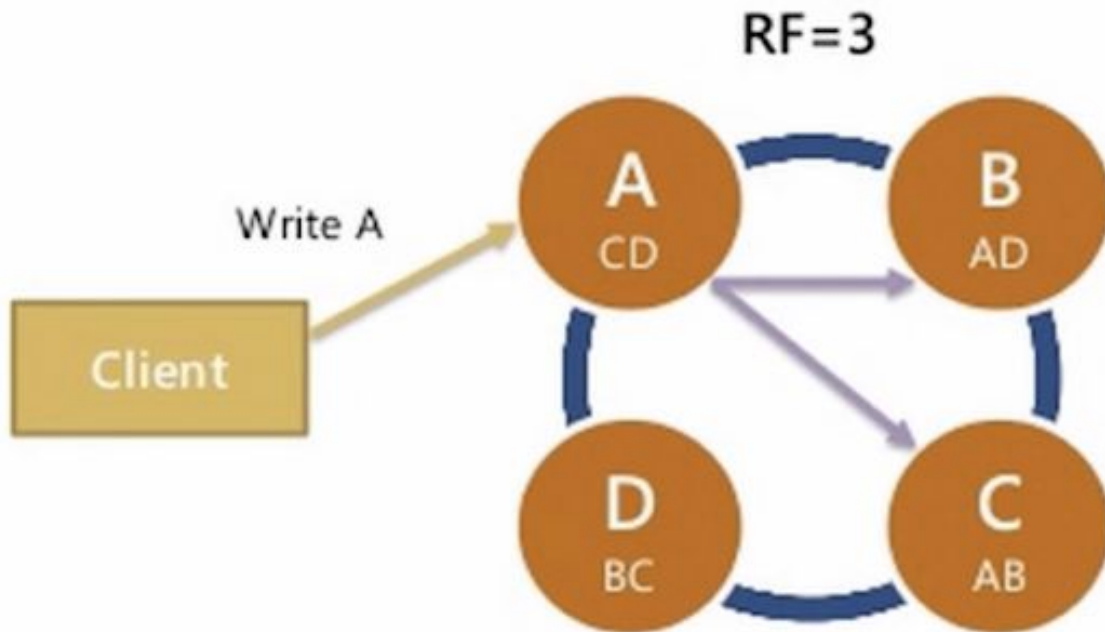


# Latency

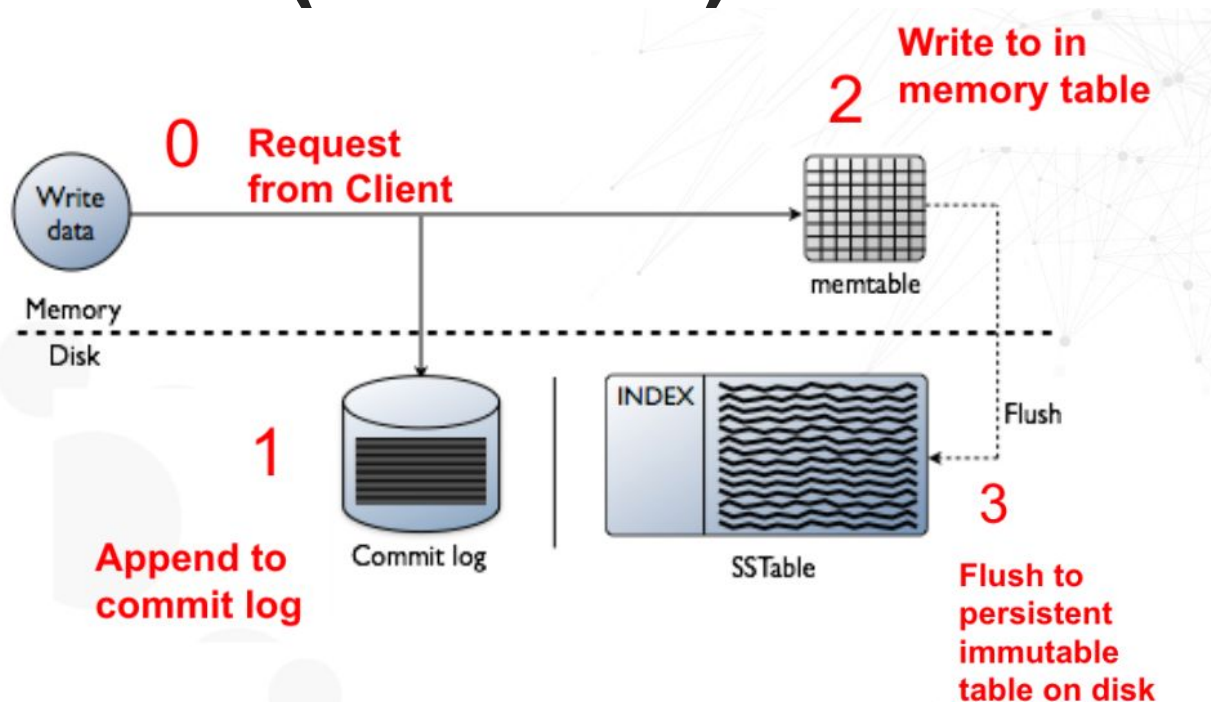
- How is low latency achieved?
- It's all about the read and write path!
  - The write path is truly beautiful in its simplicity!



# Write Path (Client to Cluster)



# Write Path (Internal)





# Read Path (Client to Cluster)

- Data modeling comes in to play here!
  - 1 “simple trick about Cassandra/Nosql”
- Partition data by nodes
- Will essentially query one node and return the data
  - Constant time READ access

```
select * from myTable where state = `CA`
```

# Data Modeling is Different

**RELATIONAL DATA MODEL != NoSQL DATA MODEL**

**CQL vs SQL**



# Data Modeling of the Past

- Normalization
  - Reduce Data Redundancy
  - Increase Data Integrity
  - 1970's
  - 3rd Normal Form
  - Natural Process



# JOINS will Save the Day --Maybe

- Combine Tables with JOINS
- Ad Hoc Queries are “okay”
- Expensive

Cartoon

id	Name	City
0110	Duck Dodgers	Burbank
0100	Marvin	Mars

```
SELECT *  
FROM cartoon  
JOIN cartoon_job  
ON  
cartoon.id =  
cartoon_job.id
```

Cartoon Job

id	Occupation
0110	Astronaut
0100	Alien

# What will Really Save the Day

- Demoralization
  - All about Performance (even in RDBMS)
  - Reducing the need for JOINS by combining/adding tables
  - Reducing time for reads
    - But will increase time for writes with data duplication

All About Cartoon

iD	Name	City	Job
0110	Duck Dodgers	Burbank	Astronaut
0100	Marvin	Mars	Alien

```
SELECT *  
FROM  
All_About_Cartoon;
```



# Data Modeling in the 21<sup>st</sup> Century

- Start with your data
- Model the data according to normal forms
  - Reduce Data Redundancy
- Determine Queries based off of this data model



# Data Modeling in the 24<sup>th</sup> 1/2 Century

- Flip it upside down!
- Start with queries/access patterns
- Create a denormalized data model
- Apply that model to your data



***Denormalization of tables in Apache Cassandra is absolutely critical***

# Distributing the Data

Primary Key

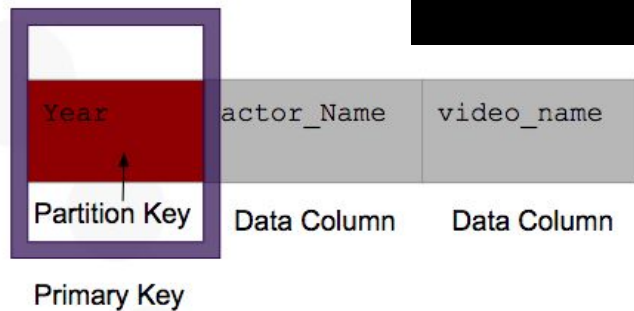
**ALL** about the **PRIMARY KEY**



# Primary Key


- Distributes the data
  - Partition Key
- Creates a Unique Key
- Made up of
  - Partition Key
  - 1 or more Clustering

```
CREATE TABLE
video_library
(year int,
actor_name text,
video_name text,
PRIMARY KEY (year))
```



# Partition Key

- Determine the distribution of data
- Partition key row value -- hashed
  - Stored on the node holds that range of values



```
INSERT INTO video_library
(year, actor_name,
video_name) VALUES (1985,
'Michael J Fox', 'Back to
the Future')
```

1985

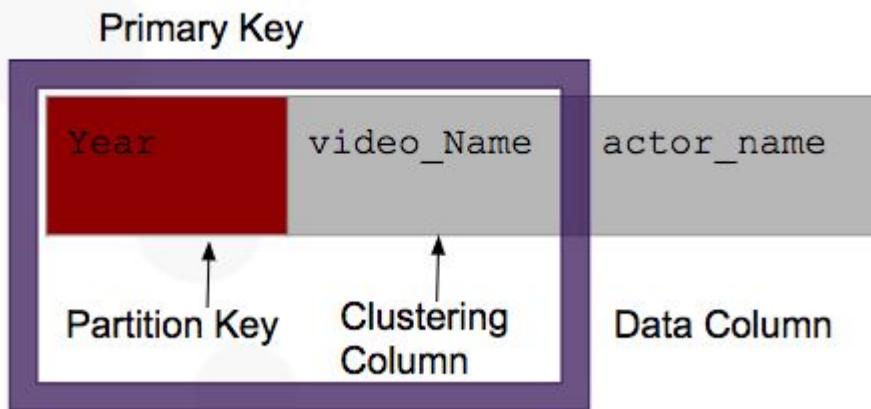
Hash  
Function

59



# Clustering Columns

- 0 or 1++ Clustering Columns
- Sort the data in sorted ascending order



```
CREATE TABLE
video_library
(year int,
actor_name text,
video_name text,
PRIMARY KEY (year),
video_name)
```

# Data Modeling Best Practices

- Do not move Relational Model **As-Is**
- 1 Table per 1 query
- Know Queries in Advance
  - Know your access patterns
- Think About your WHERE clause
- Denormalization is your Friend
  - You have money for more space, not for downtime

# Data Modeling Best Practices

- *“How to Switch from RDBMS to DynamoDB in 20 steps”*



[Jeremy Daly](#)

@jeremy\_daly

Following



STEP 5: Determine if your user access patterns require ad hoc queries that need to reshape the data. Likely the answer is, no. However, if you're building an OLAP application, [#NoSQL](#) is [#NoGOOD](#). Pat yourself on the back for trying, and use another technology.



# How to Fix Common Issues

- Data Distribution Issues
- Deleting Data Issues
- Data Integrity Issues

*Can I help you?*




# Issues with Primary Key

- Make sure Primary Key gives you unique rows
- Make sure Primary Key evenly distributes the data
  - Test and stress-test for Hot Partitions
- Make sure Partitions are not too large
  - Apache Cassandra 4.0 has a fix for this... but it's not production ready
- Until then use Bucketing to break up Large Partitions

# Tombstones

- Tombstone -- A deletion marker
  - Treats a delete as an additional insert or upsert
  - Will be removed with compaction
- When Do They Happen?
  - Deletes
  - Inserting `NULL` values
- Why so scary?
  - Impact performance -- slow reads
  - Disk space issues
  - This is why Apache Cassandra will crash!





# Tombstones --Prevent

- Avoid Inserting NULL
- Inserting a new row if data column has no value leave it blank
- When editing a row (and using the drivers) can use UNSET
- Only update fields that you need to -- don't add NULL ... add nothing!

# Batches

- Increase likelihood of having data integrity with batches
- Can put all inserts in one batch
- All will complete or an error will return after timeout

```
BEGIN BATCH;  
  
INSERT INTO video_library (year, actor_name, video_name)  
VALUES (1985, 'Michael J Fox', 'Back to the Future')  
  
INSERT INTO video_library (year, actor_name, video_name)  
VALUES  
(1985, 'Daffy', 'Duck Dodgers in the 24th and ½ Century')  
  
APPLY BATCH;
```

# That was Awesome! Now what!

- Learn more about Cassandra: <https://academy.datastax.com/>
- Learn more about DataStax: <https://www.datastax.com/>
- Follow me on Twitter: [@AmandaK\\_Data](https://twitter.com/AmandaK_Data)
- Slides will be on Github: <https://github.com/amandamoran>



O'REILLY<sup>®</sup>  
Velocity

THANK YOU



[velocityconf.com/ca](https://velocityconf.com/ca)  
#VelocityConf